

# Securing Container Runtimes

## (How Hard Could It Be?)

**Aleksa Sarai**

Senior Software Engineer

@lordcyphar

<cyphar@cyphar.com>



**OPEN** CONTAINER  
INITIATIVE

# PLEASE USE USER NAMESPACES

(Folks who did were *not vulnerable to most of these bugs...*)

# container\_runtime.pdf

- Download and extract image archive into rootfs [see my OCIv2 talk].
- Fork (and re-exec) to create proto-pid1.
  - Child will exec() pid1's code at end.
  - Parent assists during setup.

# container\_runtime (2).pdf

- Parent's job:
  - Move child process into correct cgroup.
  - Set up and container's veth or other network devices (if configured).
  - Signal child to start.

# container\_runtime (3).pdf

- Child's job:
  - Create or join namespaces (mount, pid, net, ipc, ..., and *hopefully* user).
  - Configure mountpoints for container process, pivot\_root(new root).
  - Configure seccomp filters, LSM labels, no\_new\_privs, process credentials, ...
  - Wait for parent signal, then execve(user's code).

# container\_runtime (4).pdf

- Other jobs:
  - Spawn a new process inside the container **while it's running**.
    - Rather than creating and configuring namespaces, join existing ones.
  - Modify existing container state (cgroup limits, network devices, ...).
  - Many more uninteresting things.

# CVE-2014-????

- `docker cp` didn't do *any* path sanitisation.
  - Oops.
  - `docker cp container:<symlink to /etc/shadow> w00t_w00t`
- Lesson learned:
  - Maybe we should sanitise paths.

# CVE-2015-{3627,3629,3630,3631}

- Mostly related to bad configuration or permitting bad configurations.
  - Oops.
- Lessons learned:
  - Don't make those kinds of mistakes(?).
  - **VOLUME** was probably a mistake.
  - Containers are hard.



# CVE-2016-9962

- We kept open a file descriptor to the root filesystem while joining the container.
  - Oops.
  - Container could access host through `/proc/$pid/fd/$n`.
- Lessons learned:
  - `procfs` is terrifying.
  - Make ourselves “non-dumpable” to block container process trickery.
    - Turns out there were some kernel bugs here too...

# CVE-2018-15664

- Path sanitisation isn't enough if the attacker can change the paths underneath you.
  - Oops.
  - **RENAME\_EXCHANGE** can be used to swap ("symlink-exchange") a component.
- Lessons learned:
  - Plain path sanitisation (as used to fix the 2014 bugs) is insufficient.
  - Solving this properly is non-trivial (see how LXD has done it).

# CVE-2019-5736

- We could be tricked into re-executing ourselves, pinning `/proc/self/exe`.
  - Oops.
  - This clears the “non-dumpable” bit, but maintains `/proc/self/exe`.
  - `open(“/proc/self/exe”, O_RDONLY)` then re-open it after the process dies.
- Lessons learned:
  - `procfs` is *really* terrifying.
  - Make a copy of the runc binary each time, so overwriting does nothing.
  - Maybe we should do some kernel work to block re-opens like this...

# CVE-2019-19921

- With custom images, you can use the symlink-exchange trick to mess with `/proc`.
  - Oops.
  - This means the container runtime can be tricked into not setting security labels.
  - `/proc/self/sched` can be used as a no-op writeable `procfs` target.
    - Ditto for `/proc/self/environ`.
- Lessons learned:
  - `procfs` is *absolutely horrifying*.
  - **VOLUME** (as well as several of my life decisions) were a mistake.

# What is the common theme?

- Filesystem operations are really easy to screw up.
- **p r o c f s**

**Let's make filesystem operations safe!**

# The Problem

`/foo/bar/baz`

- `baz` might be a symlink. (*Just use `O_NOFOLLOW`!*)
- `bar` might be a symlink. (*Uhhh... sanitise it in userspace?*)
- `foo` might be attacker-controlled and thus `bar` can become a symlink. (*Dammit.*)
- This *is* a solveable problem in userspace, but almost nobody does it correctly.

# openat2

```
int openat2(int dfd, const char *path,  
            struct open_how *how, size_t size);
```

```
struct open_how {  
    u64 flags;           // openat(2) flags  
    u64 mode;            // openat(2) mode  
    u64 resolve;         // RESOLVE_* flags  
    // future fields go here  
};
```



# openat2

```
#define RESOLVE_NO_SYMLINKS    ... /* Don't traverse symlinks. */  
#define RESOLVE_NO_MAGICLINKS ... /* Don't traverse magiclinks. */  
#define RESOLVE_NO_XDEV       ... /* Don't cross mounts. */  
#define RESOLVE_IN_ROOT       ... /* Resolve within a root. */
```

**libpathrs**

# Why Yet Another Library?

- Even with `openat2`:
  - Programs on old kernels still need to be hardened.
  - Users need to be **exceptionally** careful when doing other VFS operations.

# Introducing libpathrs!

- Rust library (with C bindings, usable from almost any language).
- Emulates `openat2`'s `RESOLVE_IN_ROOT` on older kernels.
- Implements helpers that match most VFS syscalls (which are correctly written).
- Includes some additional hardening (related to `procfs`).

# Usage

```
// Get a root handle for resolution.  
let root = Root::open("/path/to/root"?;  
// Resolve the path.  
let handle = root.resolve("/etc/passwd"?;  
// Upgrade the handle to a full std::fs::File.  
let file = handle.reopen(libc::O_RDONLY)?;  
  
// Or, in one line:  
let file = root.resolve("/etc/passwd"?  
                        .reopen(libc::O_RDONLY)?;
```

[docs.rs/pathrs](https://docs.rs/pathrs)

# Usage

```
root = pathrs_open("/path/to/root");
error = pathrs_error(PATHRS_ROOT, root);
if (error)
    goto err;

handle = pathrs_resolve(root, "/etc/passwd");
error = pathrs_error(PATHRS_ROOT, root);
if (error) /* or (!handle) */
    goto err;

fd = pathrs_reopen(handle, O_RDONLY);
error = pathrs_error(PATHRS_HANDLE, handle);
if (error) /* or (fd < 0) */
    goto err;

err:
if (error)
    fprintf(stderr, "Uh-oh: %s (errno=%d)\n", error->description, error->saved_errno);
pathrs_free(PATHRS_ROOT, root);
pathrs_free(PATHRS_HANDLE, handle);
pathrs_free(PATHRS_ERROR, error);
```

[docs.rs/pathrs](https://docs.rs/pathrs)

**Great! We're all done, right?**

let's have a chat about profits



# The Problem

`/proc/self/attr/exec`

- How do I make sure that I'm writing to the real `procfs` file?
  - You can grab a `/proc` handle which is definitely real (the inode is 1).
  - You can check if the target is a `procfs` file (but you aren't sure it's the right one).
  - You can disable all symlink crossings a-la `openat2` (or emulate it).
    - Wait ... how on earth do you check for bind-mounts?

year, what about binding moonlight  
to the stars?

**There is **no way** on Linux to be verify if you've crossed a bind-mount (until openat2).**

and there's magic links



/proc/self/fd/\$n  
/proc/self/exe

# YOU CAN BIND MOUNT OVER SYMLINKS

/proc/self/fd/\$n

/proc/self/exe

magic-links

**\*incomprehensible rambling\***

# Next Steps?

- Stabilise the base libpathrs C API.
- Start porting programs to libpathrs.
  - Needs Go bindings to port umoci.
- Continue kernel hardening work (which libpathrs can support opportunisically).
  - Lots of work needed to make `procfs` safe to use.



# Links

- `openat2`
  - `lwn.net/Articles/767547`
  - `lwn.net/Articles/796868`
- `libpathrs`
  - `github.com/openSUSE/libpathrs`
  - `docs.rs/pathrs`
- `github.com/cyphar/talks`

**Questions?**

# Magic-link Restriction

- Don't allow a read-only magic-link to be re-opened as read-write.
  - Requires lots of fun semantics with `O_PATH`.
  - Doesn't break userspace (based on my testing).
  - Needs to cover up a **lot** of different holes.

# O\_EMPTYPATH?

```
openat(fd, "", O_EMPTYPATH | O_RDWR);
```

# Built-in procfs handle?

```
openat(AT_PROCFD, "self/fd/$n", 0_RDWR);
```

```
setupfd = fsopen("procfs", FSOPEN_CLOEXEC);
```

```
procfd = fsmount(setupfd, FSMOUNT_CLOEXEC, 0);
```

```
openat(procfd, "self/fd/$n", 0_RDWR);
```

# Pidfd-based /proc/self ??

```
selffd = pidfd_open(getpid(), 0);  
pidfd_get_resource(selffd, PIDFD_EXE,  
                   0_RDONLY);      // ???  
openat(selffd, "exe", 0_RDONLY);    // ???
```