



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



Extensible Syscalls

Christian Brauner (Canonical)
<christian.brauner@ubuntu.com>

Aleksa Sarai (SUSE)
<cyphar@cyphar.com>



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



Groundhog Day

- We need an agreement on how we design future APIs.
 - Same argument with each new API proposal.
 - Either people want ...
 1. to have a full-blown multiplexer that is so generic that ideally it spans multiple APIs at once; or
 2. individual syscalls for each operation.
 - We've burned userspace and ourselves with both of them already.
 - There are no good guidelines to follow for developers.



LINUX
PLUMBERS
CONFERENCE

August 24-28, 2020



Groundhog Day

- We should have *documented* stronger requirements for newer syscalls.
 - A baseline of extensibility for new syscalls (to avoid the `renameat2/readlinkat/dup/accept/...` problem).
 - Note that we're not proposing that using structs in syscalls is the new baseline for extensibility. The baseline of extensibility should be the requirement of a flag argument.
 - Arguably this is the case today, though there are quite a few syscalls without a flag argument that should have one...



LINUX
PLUMBERS
CONFERENCE

August 24-28, 2020



Alternative ways of dealing with extensions

1. Always require a new syscall.
 - Costly for userspace to adapt their codebases to new syscalls all the time. This is especially true for shared libraries (doubly so with seccomp).
 - That's the other extreme of having full-blown `ioctl`-style multiplexers.
 - 64-bit flag arguments.
2. Add a new flag and have a fun time with `va_arg` (`fcntl`-style).
 - Glibc has had "fun" with this (`O_TMPFILE`), and it's pretty damn ugly.
 - *64-bit flag arguments.*
3. Fixed-size struct with "enough" padding.
 - We are bad at predicting the future.
 - Why not go the extra step and not worry about having good predictions?
4. Buy a time machine and skip extensions.
 - Can I have one too, please?
 - 128-bit flag arguments (probably).



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



Extensible structs - a modest compromise

- Introduction of new syscalls is a pain for userspace.
 - There should be a compromise between having a dedicated syscall for everything and stuffing everything into a single syscall.
- Extensible structs are a modest proposal one step further from flag-based extensibility but far away from `ioctl` madness:

`creat(2) -- open(2) -- openat2(2) -- prctl(2) -- ioctl(2)`



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



Why add more fuel to the fire?

- Support APIs that can be expected to grow beyond just adding flags.
- Need a better compromise between maintainability and simplicity.
- Alternatives are inconvenient (there are better uses for time travel).
- Better to deal with this fire here and now rather than having the same (derailing) argument on every patchset, which discourages new contributors and leads to inconsistent APIs.
 - And we can document the conclusions of this discussion to avoid bringing this issue up again.



LINUX
PLUMBERS
CONFERENCE

August 24-28, 2020



Extensible structs

```
int openat2(int fd, const char *path,  
            struct open_how *how, size_t usize);
```

- New fields always appended.
 - Zero value in new fields means “old behavior”
- **(ksize == usize)** Copy the struct verbatim.
- **(ksize > usize)** Copy usize bytes, zero-fill trailing bytes.
- **(ksize < usize)** Copy ksize bytes, check if trailing bytes are zeroed. If non-zero bytes are present, return -E2BIG.
- Dedicated kernel helper called `copy_struct_from_user()`.

```
struct open_how how = {  
    .flags      = O_RDONLY,  
    .resolve    = RESOLVE_IN_ROOT,  
    .newfield   = SOME_NEW_FLAG,  
};
```

```
int fd = openat2(dfd, path,  
                &how, sizeof(how));  
/* -E2BIG if .newfield not supported. */  
/* -EINVAL if O_RDONLY not supported. */  
if (fd < 0)  
    return -1;
```



LINUX
PLUMBERS
CONFERENCE

August 24-28, 2020



No more multiplexers

- Extensible structs are `_not_` multiplexers. They especially need to be discerned from multiplexers with polymorph types passed through a void pointer or a union or a long...
 - For example, [`bpf\(2\)`](#) is a multiplexer making use of extensible structs not an extensible struct based syscall.
 - Most people don't want to introduce any new polymorph multiplexers. That seems almost universally agreed upon.
 - glibc is explicitly advising against the introduction of new multiplexers as they are hard to deal with in shared libraries.



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



`va_arg` ... more like `bad_arg`

- `va_arg` style extension (as used by multiplexers like `fcntl(2)` and `prctl(2)`) deserve explicit mention:
 - Ugly to use due to the need to indicate which arguments should be ignored (either through flags or command enum).
 - Are limited to 5-6 arguments depending on architecture.
 - Glibc has to call `va_arg(3)` unconditionally because they don't know if a future command will require more arguments (the upshot is that they pass garbage from the stack and hope the kernel ignores it and there's enough data in the stack).



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



The “cr*p insertion vector” argument

- Someone raised the concern that extensible structs can be abused to “sneak in” problematic features without sufficient review, while being forced to introduce a new syscall in order to support the same new feature would cause a more thorough review.



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



The “cr*p insertion vector” argument

- There are at least four counter-points:
 1. For a start, the features that prompted that reaction were all caught during review so there's doubt that this is even a real problem. Especially for high-profile subsystems such as the vfs and core kernel.
 2. To the extent that this is a problem, it is no more of a problem than with flags and other extension designs -- subsystem maintainers should already be reviewing ABI changes enough to avoid this. All this design adds is the ability to easily add new fields to structs...



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



The “cr*p insertion vector” argument

3. It is debatable that syscalls prompt a more thorough review. We have a bunch of syscalls with very questionable APIs (`keyctl(2)`, `dup(2)`, etc.). [There's even an LWN article about it.](#)
4. All of the existing “questionable features” proposed could be applied to flags (`O_MAYEXEC` was originally an `openat(2)` patch). Is the ability to easily add new fields to existing struct arguments really the only thing stopping questionable features from being snuck in?



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



Checking for supported features

- Userspace needs to know what features are supported in a given syscall. This is usually done by having a (painful) trial and error approach.
 - Userspace has to write elaborate cosplay scenarios where we have to exercise the feature to see if it works without borking the system.
 - Emulation makes this worse because emulated syscalls act strange.
 - This is maybe fine for long-running system daemons but it certainly isn't workable for shared libraries or shorter-running programs (nor for our sanity -- [see the LXD code for this](#)).
- It should be possible to do this in a much simpler way -- and thanks to the design of extensible struct argument syscalls we can do it in a manner which is backward- and forward-compatible.



LINUX
PLUMBERS
CONFERENCE

August 24-28, 2020



Checking for supported features

- Here's an initial proposal.
- Already discussed on libc-alpha mailing list in the context of `clone3(2)` and [more broadly at Linux.conf.au 2020](https://linux.conf.au/2020).
- Syscall is **no-op** (returns some `errno`) and "returns" version of struct where:
 - All valid flag bits in flag fields are set.
 - All non-flag fields are filled with `0xFF`. (Could be implemented as way of describing limits for fields.)
- Userspace then looks at the final struct to determine feature support.

```
struct clone_args arg = {
    // (1 << 63)
    .flags = SUPPORTED_BITS
};
int ret = clone3(&arg, sizeof(arg));
assert(ret < 0); // always fails
if (errno != E_SUPPORTED_BITS_NOOP) {
    // kernel without feature check
    if (errno == EINVAL)
        args = CLONE3_INIT_VERSION;
    else
        return -1;
}

bool abc_supported = args.flags & CLONE_ABC;
bool xyz_supported = args.xyz != 0;
```



LINUX
PLUMBERS
CONFERENCE

August 24-28, 2020



Checking for supported features

- **(ksize == usize)** Copy the struct verbatim.
- **(ksize > usize)** Copy usize bytes, any trailing bytes on the kernel side are **ignored**.
- **(ksize < usize)** Copy ksize bytes, zero-fill the trailing (usize - ksize) bytes (unknown extensions have their zero value defined as the “old behaviour”).
- New `copy_struct_to_user()` helper.

```
struct clone_args arg = {  
    // (1 << 63)  
    .flags = SUPPORTED_BITS  
};  
int ret = clone3(&arg, sizeof(arg));  
assert(ret < 0); // always fails  
if (errno != E_SUPPORTED_BITS_NOOP) {  
    // kernel without feature check  
    if (errno == EINVAL)  
        args = CLONE3_INIT_VERSION;  
    else  
        return -1;  
}
```

```
bool abc_supported = args.flags & CLONE_ABC;  
bool xyz_supported = args.xyz != 0;
```



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



Making this part of the kernel docs

- Documentation/process/adding-syscalls.rst is of questionable currency.
 - [Our first attempt to update it](#) went largely unnoticed.
- Proposal to update adding-syscalls.rst:
 - Formalise how extensible structs are intended to be used.
 - `perf_event_open` is currently mentioned
 - ... but `clone3/openat2` are more modern.
 - Mention new variable argument or type polymorph multiplexers are not encouraged.
 - Mention syscalls with a flag argument for extensibility should use `unsigned int`.
 - Describe common issues the glibc folks have (such as pointer types in structs, et al) as well as “C annoyances” such as explicit padding and aligned fields.

commit 99c55f7d47c0dc6fc64729f37bf435abf43f4c60

Author: Alexei Starovoitov <ast@kernel.org>

Date: Fri Sep 26 00:16:57 2014 -0700

bpf: introduce BPF syscall and maps

BPF syscall is a multiplexor for a range of different operations on eBPF.
This patch introduces syscall with single command to create a map.
Next patch adds commands to access maps.

'maps' is a generic storage of different types for sharing data between kernel and userspace.

Userspace example:

```
/* this syscall wrapper creates a map with given type and attributes
 * and returns map_fd on success.
 * use close(map_fd) to delete the map
 */
```

```
int bpf_create_map(enum bpf_map_type map_type, int key_size,
                  int value_size, int max_entries)
```

```
{
    union bpf_attr attr = {
        .map_type = map_type,
        .key_size = key_size,
        .value_size = value_size,
        .max_entries = max_entries
    };
```

```
    return bpf(BPF_MAP_CREATE, &attr, sizeof(attr));
}
```

'union bpf_attr' is backwards compatible with future extensions.

More details in Documentation/networking/filter.txt and in manpage

Signed-off-by: Alexei Starovoitov <ast@plumgrid.com>

Signed-off-by: David S. Miller <davem@davemloft.net>



**LINUX
PLUMBERS
CONFERENCE**

August 24-28, 2020



`va_arg ...` more like `really_bad_arg`

- `va_arg` style multiplexers cause even more issues for glibc in general:
 - Different argument types cause confusion.
 - `size_t` and `long` cannot be differentiated and this causes issues on x32-compat due to differences between the kernel and C calling conventions (kernel requires zero or sign extension).
 - Changing a type of an argument for a single variant requires us to port all users of the syscall to a new version (see `futex(2)` and 64-bit `time_t` -- all `futex` users needed to be ported even though only a few `futex` operations actually use `time_t`).