

Operating System Security:

Attacking Container Runtimes

Aleksa Sarai

Senior Software Engineer [SUSE]

<cyphar@cyphar.com>



The Plan

- A brief introduction to containers (20').
- Filesystems and path resolution attacks (20').
- Other miscellaneous container runtime attacks (10').
- Q & A.

Containers?


- Containers are a secure^[citation needed] way of running multiple *isolated* services on a single machine without incurring the performance and resource overhead of VMs.

Virtual machines → *hardware virtualisation* → the *hardware* is a **lie**.


Containers → *operating system virtualisation* → the *OS environment* is a **lie**.

- Widely adopted in the past few years (Docker, Kubernetes, et al).

Virtual Machines




Containers




tl;dr in meme form

Virtual Machines



Containers





A Brief History Lesson.

- 1967: IBM CP-40 is “first real VM”
- 1972: IBM System/370 can host VMs.
Required additional support for hardware accelerated virtualisation.
- 2000s: Intel VT-x ('05), AMD-V ('06).
First “full virtualisation” support for x86 by adding instructions to “fake” ring-0.

- 1960s: Time-sharing systems.
- 1970s: Multi-user systems (Unix).
chroot(2) ('79).
- 2000s: Introduction of containers.
FreeBSD Jails ('00), Linux vServer ('01), Solaris Zones ('05), OpenVZ ('05), LXC ('08), Docker ('13).

Interesting talk on (part of) this history: <https://youtu.be/hgN8pCMLI2U>

chroot

- First “container” concept – introduced in Unix v7.
- Run a normal program, but pretend a host directory is the root ("/").
 - (In theory) service cannot access anything outside "/".
 - Operating system virtualises (lies about) the filesystem layout.
- Widely used for running file sharing and web servers for decades.

chroot

[good_program.gif]

Yo, can you pass me a handle to
/var/www/htdocs/index.php?

Sure, here you go!



Host Filesystem

```
/  
etc/  
  shadow  
var/run/  
  bank-records.csv  
home/pedram/  
  test-answers.txt  
var/www/htdocs/  
  index.php  
  forum.php
```

chroot

[evil_program.gif]

Yo, can you pass me a handle to
/home/pedram/test-answers.txt?

Sure, here you go!



Host Filesystem

```
/
etc/
shadow
var/run/
bank-records.csv
home/pedram/
test-answers.txt
var/www/htdocs/
index.php
forum.php
```

chroot

[good_program.gif]

Yo, can you pass me a handle to
/var/www/htdocs/index.php?

Sure, here you go!



Host Filesystem

```
/
home/pedram/
  test-answers.txt

chroot/
bin/
  good_program
  evil_program
var/www/htdocs/
  index.php
```

chroot

[evil_program.gif]

Yo, can you pass me a handle to
/home/pedram/test-answers.txt?

No idea what you're on about.



Host Filesystem

```
/
home/pedram/
  test-answers.txt
```

```
chroot/
bin/
  good_program
  evil_program
var/www/htdocs/
  index.php
```

chroot

- Sounds good, isn't that enough?

chroot

[evil_program.gif]

Yo, can I listen to packets on
0.0.0.0:443?

Sure, go right ahead!



[evil_program.gif]

Yo, can I connect to
<https://secret-network.local:1337/>?

Sure, go right ahead!



chroot

[evil_program.gif]

Yo, can I send SIGKILL
[security_policy_daemon]?

Sure, go right ahead!



[evil_program.gif]

Yo, can I fork-bomb you?

Sure, go right ahe---[crash]



Containers

- Conceptually very similar to **chroot(2)**, except now the kernel lies about:

Whether any other processes are on the system.	[PID Namespace]
What network resources are available.	[Network Namespace]
What privileges the contained program has.	[User Namespace]
...	[Time, IPC, UTS, ...]
- And resource limits can be applied to these processes (akin to Unix's ulimits).

No more fork-bombs or memory exhaustion tricks.	[cgroups]
---	-----------

Container Runtimes

- Unlike other systems (FreeBSD Jails, Solaris Zones, ...) Linux Containers are a Rube Goldberg machine of security features and isolation primitives.
 - The operating system provides individual knobs to configure them, but the final configuration is entirely up to userspace.
- Container runtimes (runc, LXC, ...) perform this and similar management tasks.
 - Tools like Docker are built on top of lower-level tools like runc.
 - They (generally speaking) *run as root* and *interact with container resources*.

Attacking the Container Runtime

Filesystem Race Conditions

Container Filesystem

- Container filesystems are just regular directories – think back to **chroot**(2).
(Though these days we actually use **pivot_root**(2).)
- Container runtimes have to do lots of operations on the container filesystem as root.
Obvious examples are **docker cp**, but many other implicit operations too.
Remember the container has *write access to its filesystem*!
Confused deputy attacks.

Obvious Attacks

```
$ docker cp container:../../../../etc/shadow shadow
$ docker cp container:symlink_to_etc_shad shadow
```

Filesystem Races


- We added some safe resolution checks to Docker in 2014, but they had a flaw:
Time of Check, Time of Use (TOCTOU).
We sanitised the path before opening it ...
... but there's a race window between sanitisation and opening.
- This is actually a fairly common Unix bug which is quite hard to get right.

The Problem

`/foo/bar/baz`

- `baz` might be a symlink. (*Just use `O_NOFOLLOW`!*)
- `bar` might be a symlink. (*Uhhh... sanitise it in userspace?*)
- `foo` might be attacker-controlled and thus `bar` can become a symlink. (*Dammit.*)
- This *is* a solveable problem, but almost nobody does it correctly.

CVE-2018-15664



```
docker cp  
ctr: /foo/bar/shadow shadow
```


Solution

- Always grab a stable handle (a file handle) to the target path and then check *that*.
If the path moves, the file handle moves with it.
- Unfortunately we haven't fully implemented this (it's quite hard to get right).

Attacking the Container Runtime

Some Other Fun Issues


procfs

- `/proc` is a very scary and magical Linux filesystem.


It is an kernel API which has a lot of holes in it.

Quite a few container breakouts were discovered by (ab)using `procfs`.

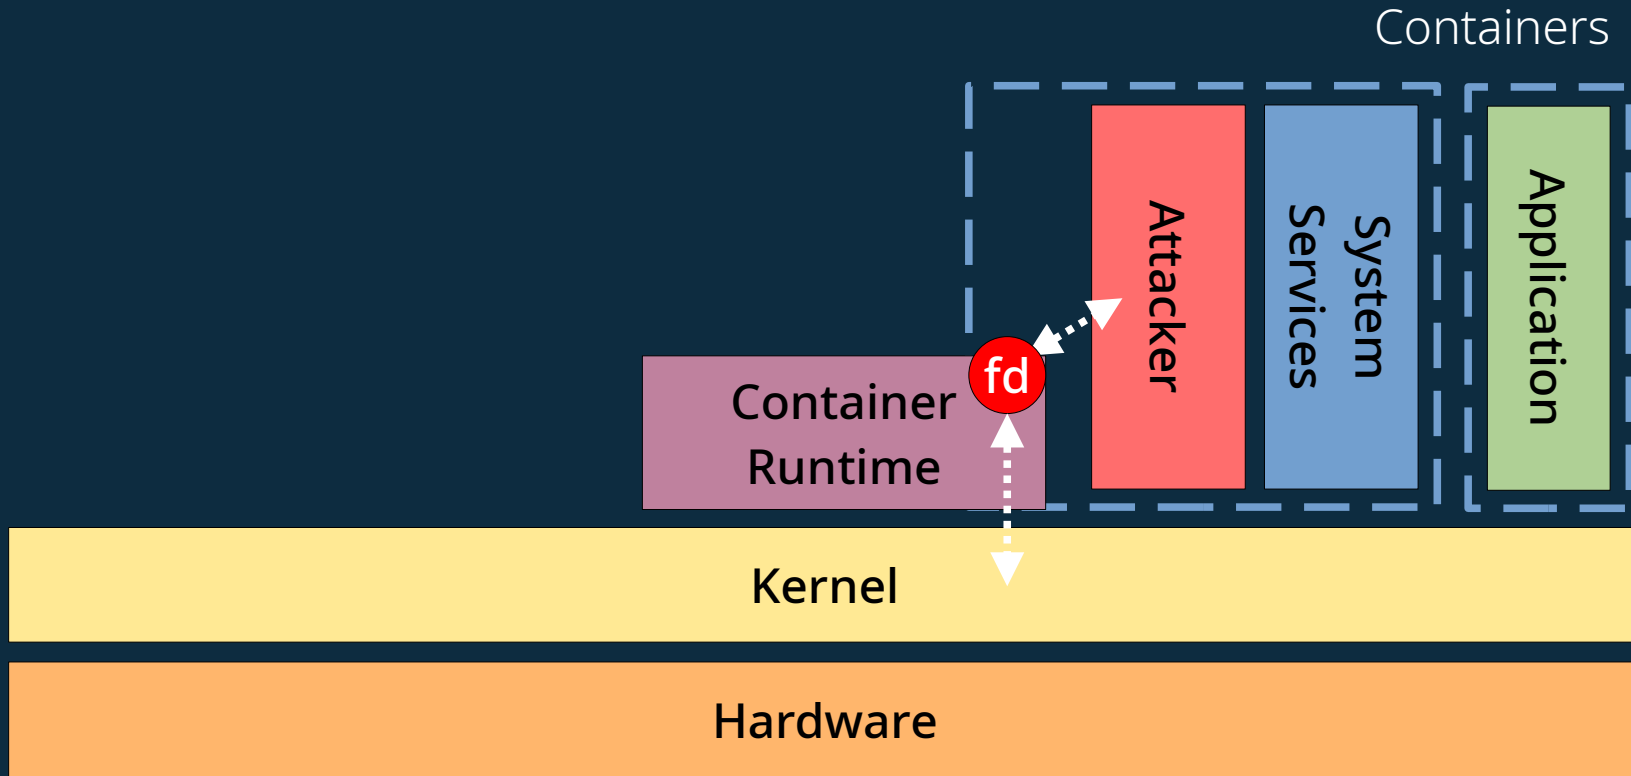
The Kernel Is Shared



The Kernel Is Shared



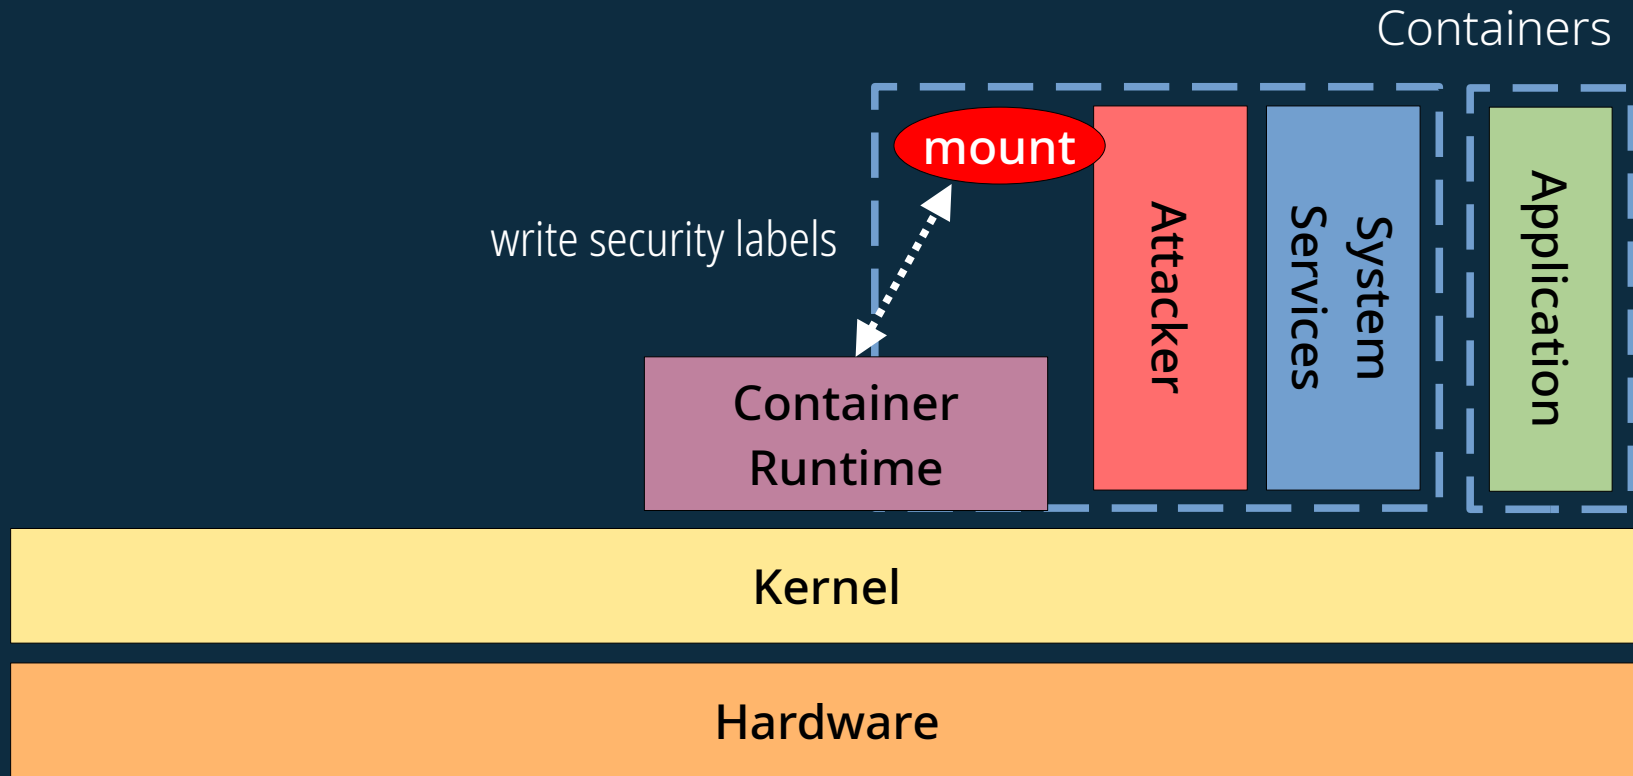
The Kernel Is Shared



CVE-2016-9962

- We kept open a file descriptor to the root filesystem while joining the container.
 - Container could access host through `/proc/$pid/fd/$n`.
- Lessons learned:
 - Make ourselves “non-dumpable” to block container process trickery.
 - Turns out there were some kernel bugs here too...

The Kernel Is Shared



CVE-2019-19921

- With custom images, you can use the symlink-exchange trick to mess with `/proc`.
 - This means the container runtime can be tricked into not setting security labels.
 - `/proc/self/sched` can be used as a no-op writeable `procfs` target.
 - Ditto for `/proc/self/environ`.
- Lessons learned:
 - `procfs` is like staring into a bottomless abyss, filled with pain and CVEs.
 - **VOLUME** were a mistake, as were several of my life decisions at this point.

Any Questions?

... or I can discuss some defensive measures.

Slides are available at github.com/cyphar/talks.

let's make filesystem operations safe!

the (old) solution

/foo/bar/baz

- For each component:
 - Open the next component (with `O_NOFOLLOW`) relative to the current one.
 - Handle symlinks in userspace by keeping track of the “text” path.
 - Do some double-checking along the way through */proc* and hope it works.
- Very hard to get right, and it looks like nobody is actually doing it.

the new solution

```
int openat2(int dfd, const char *path,  
            struct open_how *how, size_t size);
```

```
struct open_how {  
    u64 flags;           // openat(2) flags  
    u64 mode;           // openat(2) mode  
    u64 resolve;        // RESOLVE_* flags  
    // future fields go here  
};
```

openat2

```
#define RESOLVE_NO_SYMLINKS    ... /* Don't traverse symlinks. */  
#define RESOLVE_NO_MAGICLINKS ... /* Don't traverse magiclinks. */  
#define RESOLVE_NO_XDEV       ... /* Don't cross mounts. */  
#define RESOLVE_IN_ROOT       ... /* Resolve within a root. */
```

so, are we done?

- Not by a long shot.
- It's hard to get this stuff right, and even with **openat2**:
 - Programs on old kernels still need to be hardened.
 - Users need to be **exceptionally** careful when doing other VFS operations.
 - Programs need to be restructured to use file descriptors everywhere.

a library to make path resolution safe.

lib

path r

s

libpathrs

libpathrs

(a **library** to make **path** resolution **safe**.)

libpathrs

(a **library** to make **path** resolution **safe**.)

(it's also written in rust.)

introducing libpathrs!

- Rust library (with C bindings, usable from almost any language).
- Emulates `openat2`'s `RESOLVE_IN_ROOT` on older kernels.
- Implements helpers that match most VFS syscalls (which are correctly written).
- Includes some additional hardening (related to `procfs`).

usage

```
// Get a root handle for resolution.  
let root = Root::open("/path/to/root"?;  
// Resolve the path.  
let handle = root.resolve("/etc/passwd"?;  
// Upgrade the handle to a full std::fs::File.  
let file = handle.reopen(libc::O_RDONLY)?;  
  
// Or, in one line:  
let file = root.resolve("/etc/passwd"?  
                        .reopen(libc::O_RDONLY)?;
```

docs.rs/pathrs

usage

```
root = pathrs_open("/path/to/root");
error = pathrs_error(PATHRS_ROOT, root);
if (error)
    goto err;

handle = pathrs_resolve(root, "/etc/passwd");
error = pathrs_error(PATHRS_ROOT, root);
if (error) /* or (!handle) */
    goto err;

fd = pathrs_reopen(handle, O_RDONLY);
error = pathrs_error(PATHRS_HANDLE, handle);
if (error) /* or (fd < 0) */
    goto err;

err:
if (error)
    fprintf(stderr, "Uh-oh: %s (errno=%d)\n", error->description, error->saved_errno);
pathrs_free(PATHRS_ROOT, root);
pathrs_free(PATHRS_HANDLE, handle);
pathrs_free(PATHRS_ERROR, error);
```

docs.rs/pathrs

demo time.

great!
now we're all done, right?

let's have a chat about profits

the other problem

`/proc/self/attr/exec`

- How do I make sure that I'm writing to the real `procfs` file?
 - You can grab a `/proc` handle which is definitely real (the inode is 1).
 - You can check if the target is a `procfs` file (but you aren't sure it's the right one).
 - You can disable all symlink crossings a-la `openat2` (or emulate it).
 - Wait ... how on earth do you check for bind-mounts?

year, what about binding moonlight
to the sun?

There is **no way on Linux to be verify if you've crossed a bind-mount (until openat2).**

and there's magic links



/proc/self/fd/\$n
/proc/self/exe

YOU CAN BIND MOUNT OVER SYMLINKS

/proc/self/fd/\$n

/proc/self/exe

magic-links

incomprehensible rambling

next steps

- Stabilise the base libpathrs C API.
- Start porting programs to libpathrs.
- Continue kernel hardening work (which libpathrs can support opportunisically).
 - Lots of work needed to make **procfs** safe to use.

links

- `openat2` (in Linux 5.6)
 - `lwn.net/Articles/767547`
 - `lwn.net/Articles/796868`
 - `man 2 openat2`
- `libpathrs`
 - `github.com/openSUSE/libpathrs`
 - `docs.rs/pathrs`
- `github.com/cyphar/talks`

questions?

magic-link restriction

- Don't allow a read-only magic-link to be re-opened as read-write.
 - Requires lots of fun semantics with `O_PATH`.
 - Doesn't break userspace (based on my testing).
 - Needs to cover up a **lot** of different holes.

O_EMPTYPATH?

```
openat(fd, "", O_EMPTYPATH | O_RDWR);
```

built-in procfs handle?

```
openat(AT_PROCFD, "self/fd/$n", 0_RDWR);
```

```
setupfd = fsopen("procfs", FSOPEN_CLOEXEC);
```

```
procfd = fsmount(setupfd, FSMOUNT_CLOEXEC, 0);
```

```
openat(procfd, "self/fd/$n", 0_RDWR);
```

pidfd-based /proc/self ??

```
selffd = pidfd_open(getpid(), 0);  
pidfd_get_resource(selffd, PIDFD_EXE,  
                   0_RDONLY);      // ???  
openat(selffd, "exe", 0_RDONLY);    // ???
```