

---

# Path Safety

## “in the Trenches”

### 実地に踏んだパス安全の経験

CC-BY-SA 4.0



Aleksa Sarai  
[github.com/cyphar](https://github.com/cyphar)  
2025年12月9日

---

懐かしいなあ…

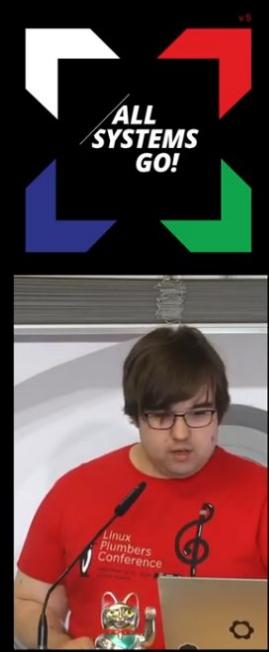




runc

## libpathrs

- Rust library that wraps the most commonly needed filesystem operations on a root filesystem with friendly™ C FFI interfaces.
  - Also has Go and Python bindings.
- Currently intended for RESOLVE\_IN\_ROOT users, but RESOLVE\_BENEATH could easily be added.
- Newer kernel features are automatically used if available.



The  
foundational  
user-space  
Linux conference

Berlin  
Sept 25-26 2024

supported by  
CHORUS ONE  
Meta

libpathrs: securing path operations for system tools



All Systems Go!

チャンネル登録  
チャンネル登録者数 4290人

チャンネル登録



Message-ID: <2025-11-05-remember-remember-the-fifth-of-november-3EtRdS@cyphar.com>  
Date: Wed, 5 Nov 2025 20:53:08 +1100  
From: Alekxa Sarai <cyphar@...har.com>  
To: oss-security@...ts.openwall.com, fulldisclosure@...lists.org  
Subject: runc container breakouts via procfs writes: CVE-2025-31133,  
CVE-2025-52565, and CVE-2025-52881

Hello,

This is a notification to vendors that use or ship runc about THREE (3) high-severity vulnerabilities (CVE-2025-31133, CVE-2025-52565, and CVE-2025-52881). All three vulnerabilities ultimately allow (through different methods) for full container breakouts by bypassing runc's restrictions for writing to arbitrary /proc files.

Message-ID: <20240903.014649-personal.smudges.long.champ-QiEEimlh1P@cyphar.com>  
Date: Tue, 3 Sep 2024 12:05:05 +1000  
From: Aleksa Sarai <cyphar@...har.com>  
To: security-announce@...ncontainers.org, oss-security@...ts.openwall.com  
Subject: CVE-2024-45310: runc can be tricked into creating empty  
files/directories on host

>  
Due to the low severity of this CVE, this security patch is being released with  
> NO embargo period.

[ Summary ]

re runc 1.1.13 and earlier as well as 1.2.0-rc2 and earlier can be tricked into  
creating empty files or directories in arbitrary locations in the host  
filesystem by sharing a volume between two containers and exploiting a race  
with os.MkdirAll. While this can be used to create empty files, existing  
files \*\*will not\*\* be truncated.

long.champ-QiEEimlh1P@cyphar.com>

Message-ID: <20210519100013.7qu6n5xtqwezmq4e@yavin>  
Date: Wed, 19 May 2021 20:00:33 +1000  
From: Alekxa Sarai <cyphar@...har.com>  
To: oss-security@...ts.openwall.com  
Subject: CVE-2021-30465: runc <1.0.0-rc95 vulnerable to symlink-exchange  
attack

This vulnerability was made public on 2021-05-19 10:00:00 UTC.

[ Summary ]

runc 1.0.0-rc94 and earlier are vulnerable to a symlink exchange attack whereby an attacker can request a seemingly-innocuous container configuration that actually results in the host filesystem being bind-mounted into the container (allowing for a container escape). CVE-2021-30465 has been assigned for this issue.

creates  
filesystem by ...  
with os.MkdirAll. While ...  
files \*\*will not\*\* be truncated.

Message-ID: <20210519100013.7qu6n5xtqwezmg4e@yau...>  
Date: Wed, 19 May 2021 20:00:33 +1000  
From: Aleksei Sarai <aleksei.sarai+1000@yau...>  
To: long.champ-QiEEimlh1P@cyphar.com>

# CVE-2019-16884: AppArmor can be bypassed by a malicious image that specifies a volume at /proc #2128

 Closed

#2128



leoluk opened on Sep 22, 2019 · edited by leoluk

Edits · ...

A malicious volume can specify a volume mount on /proc . Since Docker populates the volume by copying data present in the image, it's possible to build a fake structure that will trick runc into believing it had successfully written to /proc/self/attr/exec :

create a volume mount on /proc with a path like /proc/0/attr/exec . This will be copied into the container filesystem by the host's /proc . The file descriptor 0/465 has been assigned for this with os.MkdirAll. While files \*\*will\*\* be truncated.

shamp-QiEEimlh1P@cyphar.com>

## CVE-2018-15664: docker (all versions) is vulnerable to a symlink-race attack

From: Alekxa Sarai <cyphar () cyphar com>  
Date: Tue, 28 May 2019 14:25:13 +1000

There is no released Docker version with a fix for this issue at the time of writing. I've submitted a patch upstream[1] which is still undergoing code review, and after discussion with them they agreed that public disclosure of the issue was reasonable. Since the SUSE bug report contains exploit scripts[2], I've attached them here too.

This attack was discovered by myself (Alekxa Sarai), though Tõnis Tiigi did mention the possibility of an attack like this in the past (at the time we thought the race window was to small to exploit). In addition, you could see this exploit as a continuation of some 'docker cp' security bugs that I helped find and fix more than 4 years ago in 2014[3,4] (these were never assigned CVEs because at the time it was thought that attacks which used access to docker.sock were not valid security bugs).

[[ Overview ]]

The basic premise of this attack is that FollowSymlinkInScope suffers from a fairly fundamental TOCTOU attack. The purpose of FollowSymlinkInScope is to take a given path and safely resolve it as though the process was inside the container. After the full path has been resolved, the resolved path is passed around a bit and then operated on a bit later (in the case of 'docker cp' it is opened when creating the archive that is streamed to the client). If an attacker can add a symlink component to the path **\*after\*** the resolution but **\*before\*** it is operated on, then you could end up resolving the symlink path component on the host as root. In the case of 'docker cp' this gives you read **\*and\*** write access to any path on the host.

files

# Docker 1.12.6 - Security Advisory

CVE-1

From: Nathan McCauley <nathan.mccauley () docker com>  
Date: Tue, 10 Jan 2017 17:58:56 -0800

From: Ale

Date: Tue

There is no time of writing undergoing a public disclosure. It contains explanations.

Docker Engine version 1.12.6 has been released to address a vulnerability and is immediately available for all supported platforms. Users are advised to upgrade existing installations of the Docker Engine and use 1.12.6 for new installations.

CVE-  
image

(Close)



[ [ Overview ] ]

The basic premise from a fairly fundamental follow-symlink issue though the process has been resolved, it operated on a bit creating the architecture add a symlink component on the host read \*and\* write a

files

Credit for this discovery goes to Aleksa Sarai from SUSE and Tõnis Tiigi from Docker.

Please send any questions to security () docker com.

===== [CVE-2016-9962] Insecure opening of file-descriptor allows privilege escalation =====

RunC allowed additional container processes via `runc exec` to be ptraced by the pid 1 of the container. This allows the main processes of the container, if running as root, to gain access to file-descriptors of these new processes during the initialization and can lead to container escapes or modification of runC state before the process is fully placed inside the container

.com>

-race attack

*PLEASE USE USER NAMESPACES*

# path safety

## 「パス安全」とは

- “**regular**” path safety  
「普通の」パス安全
- “**strict**” path safety (procfs)  
「厳しい」パス安全(特にprocfsの為)

# regular path safety

## 普通のパス安全

---

# **regular path safety**

## 普通のパス安全

- Almost all system tools need to interact with unsafe paths.
- When operating on a path, a path component might be swapped with a symlink.
- Path sanitisation is useless because it just becomes TOCTTOU.

---

# **regular path (un)safety**

## 普通のパス安全のない一例

```
open("$rootfs/foo/bar/baz", ...)
```

# regular path (un)safety

## 普通のパス安全のない一例

```
open("$rootfs/foo/bar/baz", ...)
```



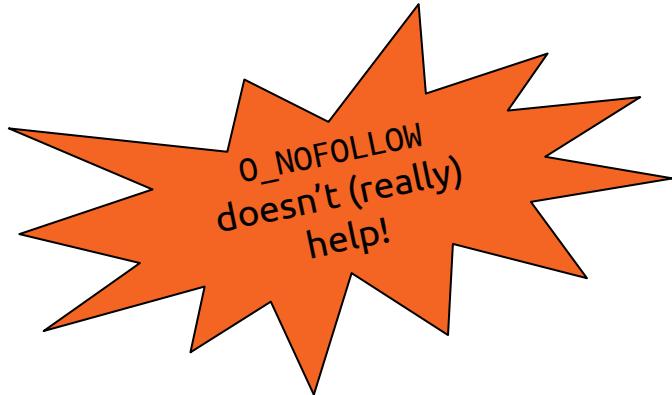
/host/path/



# regular path (un)safety

## 普通のパス安全のない一例

```
open("$rootfs/foo/bar/baz", ...)
```



/host/path/



---

# (broken) previous approaches

## 将来の駄目なやり方

### `filepath.Join`

- Trivially exploitable with `.../.../.../.../.../.../.../.../foo`.
- Symlink components are resolved *in the host*.

# example vulnerabilities (i)

## 脆弱性の一例(其ノ壱)

Access to docker daemon gives unrestricted read access  
to host filesystem #5656

 Closed  #5720

 Supermathie opened on May 8, 2014  ...

Kind of under the same umbrella, access to talk to the docker daemon allows access to the entire host filesystem:

```
○ .. ls id_dsa
ls: cannot access id_dsa: No such file or directory

○ .. docker.io cp 2227e2a5bd3f:../../../../../../../../root/.ssh/id_dsa .
2014/05/07 11:28:23 lchown id_dsa: operation not permitted

○ .. ls id_dsa
id_dsa

○ .. cat id_dsa
SECRET_KEY!
```

Tested with docker 0.9.1 and 0.10.0



# example vulnerabilities (i)

## 脆弱性の一例(其ノ壱)

Access to docker daemon gives unrestricted read access  
to host filesystem #5656

```
○ → docker.io cp 2227e2a5bd3f:../../../../../../../../root/.ssh/id_dsa .
2014/05/07 11:28:23 lchown id_dsa: operation not permitted
```

```
○ → ls id_dsa
id_dsa
```

```
○ → cat id_dsa
SECRET_KEY!
```

```
○ → cat id_dsa
SECRET_KEY!
```

Tested with docker 0.9.1 and 0.10.0



# example vulnerabilities (i)

## 脆弱性の一例(其ノ壱)

Access to docker daemon gives unrestricted read access

jpetazzo on May 8, 2014

Hi,

This is normal, and expected. Access to the Docker daemon should be protected. See #560 and [the documentation](#).

TL;DR: by default, only root can access the Docker socket (and the `docker` group, but by default, nobody is in that group).

Thank you!

jpetazzo closed this as completed on May 8, 2014

Tested with docker 0.9.1 and 0.10.0

# example vulnerabilities (i)

## 脆弱性の一例(其ノ壱)

Access to docker daemon gives unrestricted read access

jpetazzo on May 8, 2014

This is normal, and expected.

jpetazzo closed this as completed on May 8, 2014

Tested with docker 0.9.1 and 0.10.0

This is normal, and expected.



# example vulnerabilities (i)

## 脆弱性の一例(其ノ壱)

A screenshot of a GitHub pull request page. The title of the pull request is "Ensure docker cp cannot traverse outside container rootfs #5720". It was created by [jpetazzo](#) on May 14, 2014, and merged by [crosbymichael](#) on May 15, 2014. The commit message states: "Access to docker daemon gives ...". A comment from [cyphar](#) on May 10, 2014, provides details: "This patch fixes the bug (#5656) that allowed cp to copy files outside of the containers rootfs, by passing a relative path (such as ../../../../../../etc/shadow). This is fixed by first converting the path to an absolute path (relative to /) and then appending it to the container's rootfs before continuing." The GitHub interface shows 12 conversations, 3 commits, 0 checks, and 4 files changed.

Access to docker daemon gives ...

jpetazzo on May 14, 2014

Ensure docker cp cannot traverse outside container rootfs #5720

Merged

crosbymichael merged 3 commits into [moby:master](#) from [cyphar:5656-cp-absolute-paths](#) on May 15, 2014

Conversation 12 Commits 3 Checks 0 Files changed 4

Contributor ...

cyphar commented on May 10, 2014

This patch fixes the bug (#5656) that allowed cp to copy files outside of the containers rootfs, by passing a relative path (such as ../../../../../../etc/shadow). This is fixed by first converting the path to an absolute path (relative to /) and then appending it to the container's rootfs before continuing.

Docker-DCO-1.1-Signed-off-by: Aleksa Sarai [cyphar@cyphar.com](mailto:cyphar@cyphar.com) (github: cyphar)

---

## filepath-securejoin

- Path lookups in userspace with chroot(2)-like scoping.
-

---

# **still broken previous approaches**

## まだ駄目なやり方

### **filepath-securejoin**

- Path lookups in userspace with chroot(2)-like scoping.

---

# **still broken previous approaches**

## まだ駄目なやり方

### **filepath-securejoin**

- Path lookups in userspace with chroot(2)-like scoping.
- Returns a path string – not safe against races!

---

# **example vulnerabilities (ii)**

## 脆弱性の一例(其ノ式)

**CVE-2018-15664**

- docker cp used chroot(2) to try to block path-based attacks.
- Unfortunately, Docker rooted the chroot(2) within the attacker-controlled path and so the same issue was present.
- (This was one example of a larger pattern within Docker that AFAIK is still not fully resolved.)

---

# **example vulnerabilities (iii)**

## 脆弱性の一例(其ノ参)

**CVE-2024-45310**

- Mountpoint targets would be created with `os.MkdirAll` or `O_CREAT` of a `filepath-securejoin` path.
- Racing attackers could trick us into creating host inodes.

---

# solutions (i)

## 解決(其ノ壱)



### openat2(2)

(since Linux 5.6)

- chroot(2)-like IN\_ROOT Just Works™ (mostly).
- NO\_SYMLINKS, BENEATH for everything else.
- NO\_MAGICLINKS, NO\_XDEV are particularly useful.
- Requires moving to a more file-descriptor based approach.

---

# solutions (ii)

## 解決(其ノ式)



### openat(0\_PATH)

(since Linux 2.6.39-ish)

- Manually do lookup with 0\_PATH handles, emulating what openat2(2) does.
- .. and / components are usually verified through /proc/self/fd.
- Still requires moving to a more file-descriptor based approach.

---

# **example vulnerabilities (iv)**

## 脆弱性の一例(其ノ肆)

**CVE-2025-31133 & CVE-2025-52565**

- Special bind-mount sources like `/dev/null` and `/dev/pts/$n` were sourced from the container.
- Racing attackers could replace these with symlinks, causing us to bind-mount procfs files that would normally be masked, allowing for breakouts.

strict path safety

厳しいパス安全

---

# **strict path safety**

## 厳しいパス安全

- Some filesystem operations need to operate on a specific file, the most common example is operations on procfs.
- In these cases, bind-mounts and symlinks are big trouble.

---

# **strict path (un)safety**

## 厳しいパス安全のない一例

```
open("/proc/self/attr/exec", ...)  
open("/proc/sys/some/sysctl", ...)  
open("/proc/self/fd/$fd", ...)  
execve("/proc/self/exe", ...)
```

# strict path (un)safety

## 厳しいパス安全のない一例



```
open("/tmp/bad_binary", ...)  
open("/proc/self/attr/exec", ...)  
open("/proc/sys/some/sysctl", ...)  
open("/proc/self/fd/$fd", ...)  
execve("/proc/self/exe", ...)
```

The code snippet shows several system calls. Red annotations highlight specific paths:

- /tmp/bad\_binary: A red arrow points from this label to the first argument of the first open call.
- /proc/self/sched: A red arrow points from this label to the first argument of the second open call.
- /proc/sys/kernel/core\_pattern: A red arrow points from this label to the first argument of the last execve call.

---

# **example vulnerabilities (v)**

## 脆弱性の一例(其ノ伍)

### CVE-2019-19921

- runc would previously just write to /proc normally.
- Attackers could trick runc into configuring a tmpfs mount on top of /proc with fake procfs files.
- LSM labels are set through /proc/self/attr/\$lsm/exec, fake files make this a no-op.

---

# **temporary fix**

## 暫定手段

- Use `fstatfs` to check that we are operating on a procfs file.
  - Only a temporary fix – we can't be sure which procfs file it is!
-

---

# example vulnerabilities (vi)

## 脆弱性の一例(其ノ陸)

CVE-2025-5288

- Rather than using tmpfs, trick runc into bind-mounting bits of procfs.
  - /proc/self/attr/exec ← /proc/self/sched (no-op)
  - /proc/self/attr/exec ← /proc/sysrq-trigger (crash)
  - /proc/sys/... ← /proc/sys/kernel/core\_pattern 😈

---

# solutions (iii)

## 解決(其ノ参)



### `fsopen(2) · open_tree(2)`

(since Linux 5.1)

- Provides a mechanism for a private procfs mount that cannot have overmounts or racing attackers doing mounts.
- For regular procfs files, `openat2(NO_XDEV)` is usually enough.
  - (If you check that `/proc` is `PROC_ROOT_INO`.)

Message-ID: <2025-11-05-remember-remember-the-fifth-of-november-3EtRdS@cyphar.com>  
Date: Wed, 5 Nov 2025 20:53:08 +1100  
From: Alekxa Sarai <cyphar@...har.com>  
To: oss-security@...ts.openwall.com, fulldisclosure@...lists.org  
Subject: runc container breakouts via procfs writes: CVE-2025-31133,  
CVE-2025-52565, and CVE-2025-52881

Hello,

This is a notification to vendors that use or ship runc about THREE (3) high-severity vulnerabilities (CVE-2025-31133, CVE-2025-52565, and CVE-2025-52881). All three vulnerabilities ultimately allow (through different methods) for full container breakouts by bypassing runc's restrictions for writing to arbitrary /proc files.

```
commit 2f74f4ae1b483b3a4d197f574a37c7b768ba96f
Merge: fb01482d120f 3f925525b44d
Author: Aleksa Sarai <cyphar@cyphar.com>
Date:   Wed Nov 5 20:19:30 2025 +1100

    merge private security patches into opencontainers/runc:main

Alekxa Sarai (21):
    rootfs: re-allow dangling symlinks in mount targets
    openat2: improve resilience on busy systems
    selinux: use safe procfs API for labels
    rootfs: switch to fd-based handling of mountpoint targets
    libct/system: use securejoin for /proc/$pid/stat
    init: use securejoin for /proc/self/setgroups
    init: write sysctls using safe procfs API
    utils: remove unneeded EnsureProcHandle
    utils: use safe procfs for /proc/self/fd loop code
    apparmor: use safe procfs API for labels
    ci: add lint to forbid the usage of os.Create
    rootfs: avoid using os.Create for new device inodes
    internal: add wrappers for securejoin.Proc*
    go.mod: update to github.com/cyphar/filepath-securejoin@v0.5.0
    console: verify /dev/pts/ptmx before use
    console: avoid trivial symlink attacks for /dev/console
    console: add fallback for pre-TIOCGPTPEER kernels
    console: use TIOCGPTPEER when allocating peer PTY
    *: switch to safer securejoin.Reopen
    internal: move utils.MkdirAllInRoot to internal/pathrs
    internal/sys: add VerifyInode helper

Li Fubang (1):
    libct: align param type for mountCgroupV1/V2 functions

Kir Kolyshkin (3):
    libct: maskPaths: don't rely on ENOTDIR for mount
    libct: maskPaths: only ignore ENONENT on mount dest
    libct: add/use isDevNull, verifyDevNull

Fixes: CVE-2025-31133 GHSA-9493-h29p-rfm2
Fixes: CVE-2025-52565 GHSA-qw9x-cqr3-wc7r
Fixes: CVE-2025-52881 GHSA-cgrx-mc8f-2prm
Reported-by: Lei Wang <ssst0n3@gmail.com>
Reported-by: Li Fubang <lifubang@acmCoder.com>
Reported-by: Tõnis Tiigi <tonistiigi@gmail.com>
Reported-by: Aleksa Sarai <cyphar@cyphar.com>
Signed-off-by: Aleksa Sarai <cyphar@cyphar.com>
```

111 files changed, 9452 insertions(+), 1338 deletions(-)



```
111 files changed, 9452 insertions(+), 1338 deletions(-)
```



---

# our approach

## 我らの解決手段

- Switch to [pathrs-lite](#).
  - openat2(2) support (with O\_PATH fallback).
  - Provides helpers for safe procfs operations.
  - A pure-Go version of [libpathrs](#).
- Harden verification of all special inodes we interact with.
- Audit all write paths – can they write to /proc?
- Switch to file descriptors as much as possible.

---

# general takeaways

## 解決手段

- Most system tools need at least regular path safety.
- Switch to tools like `libpathrs` (or `paths-lite`) and move to a move file-descriptor-based design.
  - If you don't need to support old kernels, feel free to use `openat2(2)` directly.
- If you use procfs, you should **really** use `libpathrs`.

---

# remaining work

## まだまだ…

- Switch runc to `libpathrs`.
  - Polish `libpathrs` some more.
- Move even more things to use file descriptor based.
  - Switch to the `fsopen(2)` mount API.
- Come up with a decent threat model for runc.

# questions?

(rants, pitchforks...?)

CC-BY-SA 4.0



---

# libpathrs

- Rust library that wraps the most commonly needed filesystem operations on a root filesystem with friendly™ C FFI interfaces.
  - Also has Go and Python bindings.
- Currently intended for RESOLVE\_IN\_ROOT users, but RESOLVE\_BENEATH could easily be added.
- Newer kernel features are automatically used if available.

# libpathrs (rust)

```
use pathrs::{Root, flags::OpenFlags};
let root = Root::open("/path/to/root")?;

// Resolve and open a file.
let passwd = root
    .resolve("/etc/passwd")?
    .reopen(OpenFlags::O_RDONLY)?;
// ... or ...
let passwd = root.open_subpath(
    "/etc/passwd", OpenFlags::O_RDONLY
)?;

// Create a new file and open it (O_CREAT).
let newfile = root.create_file(
    "/etc/newfile",
    OpenFlags::O_RDWR,
    &Permissions::from_mode(0o755),
)?;

// Create a symlink.
let newfile = root.create(
    "/link",
    &InodeType::Symlink("/target".into()),
)?;

// mkdir -p
let dir = root.mkdir_all(
    "/foo/bar/baz",
    &Permissions::from_mode(0o755),
)?;

// rm -r
root.remove_all("/foo/bar")?;

// See the docs for more info.
```

# libpathrs (c)

```
int root = pathrs_open_root("/path/to/root");
if (root < 0) {
    liberr = root;
    goto err;
}
int handle = pathrs_inroot_resolve(root,
                                    "/etc/passwd");
if (handle < 0) {
    liberr = handle;
    goto err;
}
int fd = pathrs_reopen(handle, O_RDONLY);
if (fd < 0) {
    liberr = fd;
    goto err;
}
// or pathrs_inroot_open(root, "/etc/passwd", ...)

err:
if (liberr < 0) {
    pathrs_error_t *error =
        pathrs_errorinfo(liberr);
    fprintf(stderr,
            "Uh-oh: %s (errno=%d)\n",
            error->description,
            error->saved_errno);
    pathrs_errorinfo_free(error);
}
close(root);
close(handle);
/* ... do something with fd ... */
```

# “reopen”?

```
use pathrs::{Root, flags::OpenFlags};  
let root = Root::open("/path/to/root")?;  
  
// Get a reusable (O_PATH) ptmx handle.  
let ptmx = root  
    .resolve("/dev/pts/ptmx")?;  
  
// Create several new console instances.  
// They are all independent instances.  
let console1 = ptmx  
    .reopen(OpenFlags::O_RDWR)?;  
let console2 = ptmx  
    .reopen(OpenFlags::O_RDWR)?;  
let console3 = ptmx  
    .reopen(OpenFlags::O_RDWR)?;
```

- Sometimes you need to re-open the same file multiple times.
- This is not just dup! It's a proper *race-free* open.
- Implementing lookups entirely with O\_PATH and re-opening is simpler...

---

# procfs api

- Detecting attackers is the primary goal, followed by resiliency.
- Private procfs instance with `fsopen` and `open_tree` if possible.
  - Can't be used for unprivileged programs...
- `openat2` or very restrictive `O_PATH` resolver for lookups.
  - *(This resolver doesn't do reopens!)*
- Used internally by `libpathrs` to implement the main API.

# procfs api (rust)

```
use pathrs::{flags::OpenFlags, procfs::*};

// Open *regular* file.
let attr_file = ProcfsHandle::new()?
    .open(
        ProcfsBase::ProcThreadSelf,
        "attr/exec",
        OpenFlags::O_WRONLY,
    )?;

// Create your own private handle.
let handle =
    ProcfsHandleBuilder::new()
        .unmasked(true)
        .build()?;
// Open a magic-link.
let exe = ProcfsHandle::new()?.open_follow(
    ProcfsBase::ProcSelf,
    "exe",
    OpenFlags::O_RDONLY,
)?;

// Equivalent to readlinkat(fd, "").
let fd_path = ProcfsHandle::new()?.readlink(
    ProcfsBase::ProcThreadSelf,
    format!("fd/{}", file.as_raw_fd()),
    OpenFlags::O_RDONLY,
)?;
```

# procfs api (c)

```
int write_apparmor_label(const char *label)
{
    /* Open *regular* file. */
    int fd = pathrs_proc_open(
        PATHRS_PROC_THREAD_SELF,
        "attr/apparmor/exec",
        O_WRONLY|O_NOFOLLOW
    );
    if (fd < 0) {
        pathrs_error_t *e = pathrs_errorinfo(fd);
        /* ... print the error ... */
        pathrs_errorinfo_free(e);
        return -1;
    }
    int err = write(fd, label, strlen(label));
    close(fd);
    return err;
}

int get_self_exe(void)
{
    /* Follows the magic-link! */
    int fd = pathrs_proc_open(PATHRS_PROC_SELF,
                               "exe", O_PATH);
    if (fd < 0) {
        pathrs_error_t *e = pathrs_errorinfo(fd);
        /* ... print the error ... */
        pathrs_errorinfo_free(e);
        return -1;
    }
    return fd;
}
```

---

# procfs limitations

- For non-magic-links, `openat2` (Linux 5.6) is sufficient.
  - (`openat2` might be blocked due to seccomp limitations.)
- For magic-links, we need:
  - `fsopen` or `open_tree` (Linux 5.1) for race safety.

# questions?

(rants, pitchforks...?)

CC-BY-SA 4.0



⋮ ⓘ CVE-2025-52881: fd reopening  
causes issues with AppArmor profiles  
(open sysctl  
net.ipv4.ip\_unprivileged\_port\_start file:  
reopen fd 8: permission denied)

#4968 · cyphar opened last month

⋮

92

func

AppArmor

fsopen("proc") → "/"

func

AppArmor

`fsopen("proc")` → "/"

`open(procfd, "sys/...")` → "/sys/..."

func

AppArmor

`fsopen("proc")` → "/"

`open(procfd, "sys/...")` → "/sys/..."



# questions?

(rants, pitchforks...?)

CC-BY-SA 4.0



---

# kernel hardening – magic-links

- Magic-links can be used to break out of containers.
  - Userspace needs to be careful about leaks.
  - Container runtimes need to use PR\_SET\_DUMPABLE!
- [CVE-2019-5736](#): Overwrite host binary with /proc/self/exe.
  - Solution: [restrict reopening with extra permissions](#).
  - (Also allow users to specify fd restrictions with openat2.)
- [Restrict mounts on top of all magic-links](#).

---

# kernel work – openat2

- Some more things we might want to add:
  - RESOLVE\_NO\_BLOCK (NO\_REMOTE?) to avoid DoSes.
  - Restrict types of files we want to open (another DoS).
  - RESOLVE\_NO\_DOTDOT for extreme lookup restrictions.
- What about atomic mknod combined with open? (O\_MKNOD?)
  - If it could take RESOLVE\_\* flags that would be even nicer!