

Path Safety “in the Trenches”

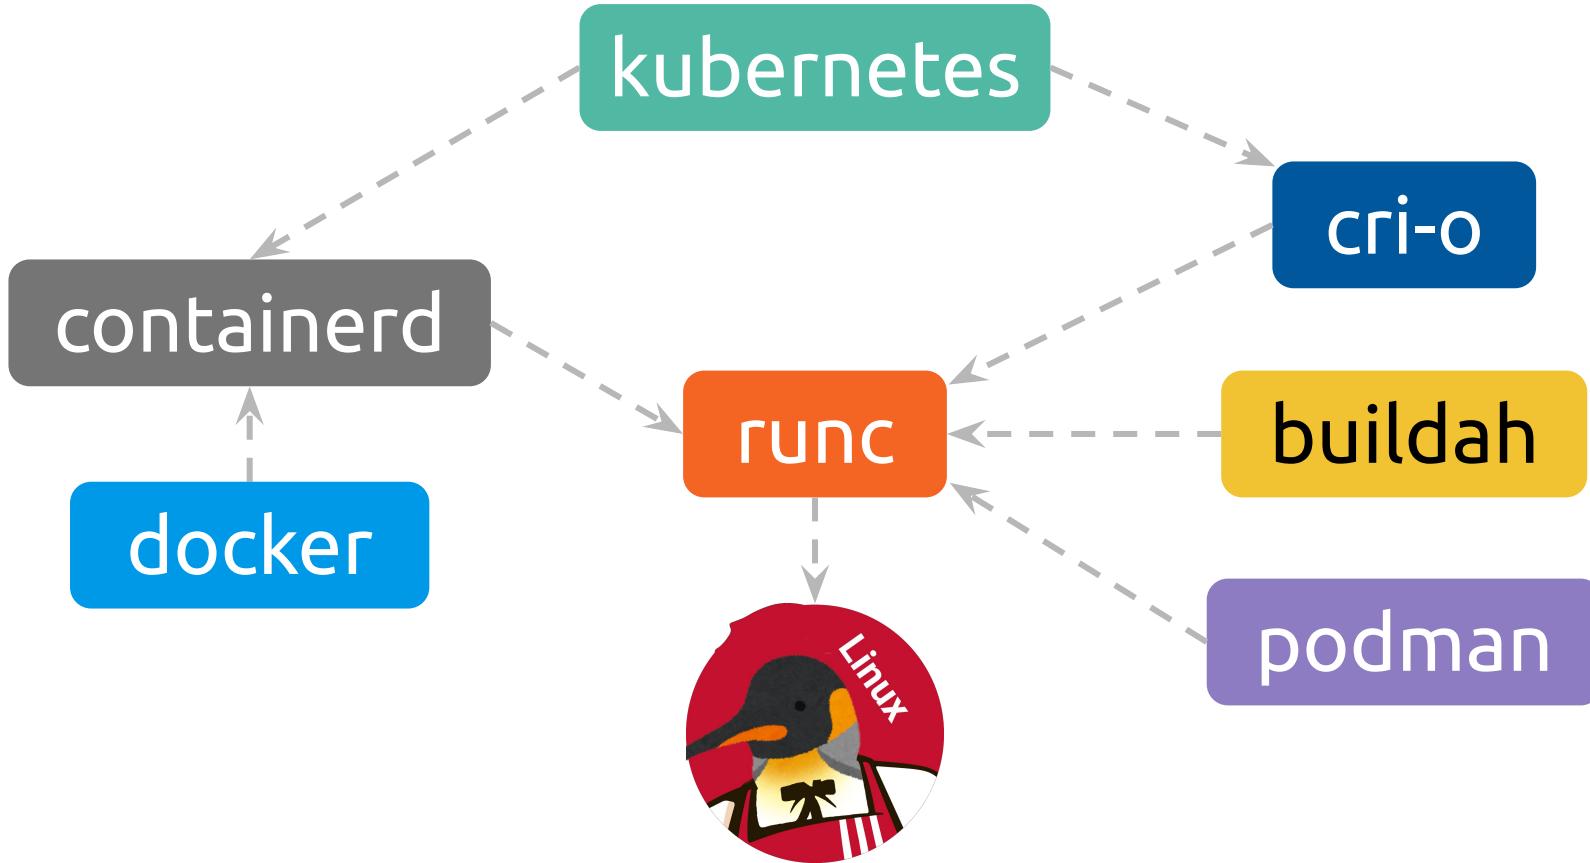
CC-BY-SA 4.0



Aleksa Sarai – github.com/cyphar
Founding Engineer @ Amutable
FOSDEM 2026-01-31



runc



```
$ cat config.json
{
    "process": { ... },
    "root": { "path": "path/to/rootfs" },
    "mounts": [
        { "destination": "/proc", "type": "proc" },
        ...
    ],
    "linux": {
        "resources": { "devices": [ ... ] },
        "namespaces": [ ... ],
        ...
    }
}
```

```
$ cat config.json
{
  ...
}
$ runc run ctr-name
ctr# ...
```

path safety

“**regular**” path safety
“**strict**” path safety (procfs)

regular path safety

- When operating on a path, a path component might be swapped with a symlink or moved.
 - Classic time-of-check-to-time-of-use attacks abound.
-

example: regular path (un)safety

```
int fd1 = open("/rootfs/etc/foo", O_CREAT|O_TRUNC|O_RDWR, 0755);
int fd2 = open("/rootfs/etc/hosts", O_NOFOLLOW|O_RDONLY);

mkdir("/rootfs/foo", 0755);
mkdir("/rootfs/foo/bar", 0755);

link("/rootfs/foo/bar/baz", "/rootfs/foo/bar/boop");

unlink("/rootfs/foo/bar/baz");
```

example: regular path (un)safety

```
int fd1 = open("/rootfs/etc/foo", O_CREAT|O_TRUNC|O_RDWR, 0755);
int fd2 = open("/rootfs/etc/hosts", O_NOFOLLOW|O_RDONLY);

mkdir("/rootfs/foo", 0755);
mkdir("/rootfs/foo/bar", 0755);

link("/rootfs/foo/bar/baz", "/rootfs/foo/bar/boop");

unlink("/rootfs/foo/bar/baz");
```

surely this isn't *that* common...

surely this isn't *that* common...

- [CVE-2017-1002101](#)
- [CVE-2018-15664](#)
- [CVE-2019-16884](#)
- [CVE-2019-19921](#)
- [CVE-2021-30465](#)
- [CVE-2023-27561](#)
- [CVE-2023-28642](#)
- [CVE-2024-1753](#)
- [CVE-2024-45310](#)
- [CVE-2024-0132](#)
- [CVE-2024-0133](#)
- [CVE-2024-9676](#)
- [CVE-2025-31133](#)
- [CVE-2025-52565](#)
- ... and so on ...

regular path safety – openat2

```
int openat2(int dirfd, const char *path,  
            struct open_how *how, size_t size);  
  
struct open_how {  
    u64 flags;      /* 0_* flags */  
    u64 mode;       /* O_CREAT file mode */  
    u64 resolve;    /* resolution flags */  
};
```

regular path safety – openat2

- Most programs can make do with one of the following:
 - RESOLVE_IN_ROOT – chroot(2)-like lookups
 - RESOLVE_BENEATH – restricted lookups
 - RESOLVE_NO_SYMLINKS – better O_NOFOLLOW (/ still escapes)
- Requires the program to primarily use file descriptors.

example: regular path safety (i)

```
int root = open("/rootfs", O_DIRECTORY|O_PATH);

struct open_how how = { .resolve = RESOLVE_IN_ROOT };

how.flags = O_CREAT|O_TRUNC|O_RDWR;
how.mode = 0755;
int fd1 = openat(root, "/etc/foo", &how, sizeof(how));

how.flags = O_NOFOLLOW|O_TRUNC|O_RDWR;
how.mode = 0;
int fd2 = openat(root, "/etc/hosts", &how, sizeof(how));
```

example: regular path safety (ii)

```
int root = open("/rootfs", O_DIRECTORY|O_PATH);

struct open_how how =
{ .flags = O_DIRECTORY|O_PATH,
  .resolve = RESOLVE_IN_ROOT };

int dfd = openat(root, "/foo", &how, sizeof(how));
mkdirat(dfd, "bar", 0755); /* mkdir(/rootfs/foo/bar) */
```

regular path safety – 0_PATH

- Implement per-component lookups in userspace (very finicky).
- Still requires file-descriptor-based code.
- Usually needs `readlink("/proc/self/fd/$n")` verification.
- *See:* systemd's `chaseat`, Go's `os.Root`, `libpathrs`.

libpathrs

- Rust library that wraps the most commonly needed filesystem operations on a root filesystem with friendly™ C FFI interfaces.
 - Also has Go and Python bindings.
 - [pathrs-lite](#) – pure-Go port.
- Transparently supports `openat2` and the `O_PATH` fallback.

example: libpathrs (c)

```
#include <pathrs.h>

int root = pathrs_open_root("/rootfs"); /* or open(2) */

int fd1 = pathrs_inroot_creat(root, "/etc/foo",
                             O_RDWR|O_TRUNC, 0644);
int fd2 = pathrs_inroot_open(root, "/etc/hosts",
                           O_NOFOLLOW|O_RDONLY);

pathrs_inroot_hardlink(root, "/foo/bar/baz", "/foo/bar/boop");

pathrs_inroot_unlink(root, "/foo/bar/baz");
```

example: libpathrs (rust)

```
use pathrs::Root, OpenFlags;
let root = Root::open("/path/to/root")?;

// Resolve and open a file.
let passwd = root
    .resolve("/etc/passwd")?
    .reopen(OpenFlags::O_RDONLY)?;
// ... or ...
let passwd = root.open_subpath(
    "/etc/passwd", OpenFlags::O_RDONLY
)?;

// Create a new file and open it (O_CREAT).
let newfile = root.create_file(
    "/etc/newfile",
    OpenFlags::O_RDWR,
    &Permissions::from_mode(0o755),
)?;

// Create a symlink.
let newfile = root.create(
    "/link",
    &InodeType::Symlink("/target".into()),
)?;

// mkdir -p
let dir = root.mkdir_all(
    "/foo/bar/baz",
    &Permissions::from_mode(0o755),
)?;

// rm -r
root.remove_all("/foo/bar")?;

// See the docs for more info.
```

-

strict path safety

`/proc/self/attr/exec,`
`/proc/self/mountinfo,`
`/proc/sys/$sysctl, etc.`

strict path safety

- For certain pseudo-filesystems we need to ensure we are operating on an exact path.
 - `procfs` is most critical.
 - Overmounts or fake mounts can trick us into doing dangerous operations or make operations a no-op.
-

example: strict path (un)safety

```
int lfd = open("/proc/self/attr/exec", O_RDWR);
dprintf(lfd, "exec docker-default\n");

int pfd = open("/proc/sys/net/ipv4/ping_group_range", O_RDWR);
dprintf(pfd, "0 0\n");

int reopen = open("/proc/thread-self/fd/123", O_RDWR);

execve("/proc/self/exe", ...);
```

surely not!

okay, it's a *bit* less common...

- [CVE-2019-16884](#)
- [CVE-2019-19921](#)
- [CVE-2025-52881](#)

example: strict path safety

```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{
    .flags = O_WRONLY,
    .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int lfd = openat(procfd, "self/attr/exec", &how, sizeof(how));
dprintf(lfd, "exec docker-default\n");
```

example: strict path safety

```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{
    .flags = O_WRONLY,
    .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int lfd = openat(procfd, "self/attr/exec", &how, sizeof(how));
dprintf(lfd, "exec docker-default\n");
```



-

magic-links?

`/proc/self/exe,`

`/proc/self/fd/$n,`

`/proc/self/root`, etc.

example: strict path safety

```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{ .flags = O_RDWR,
  .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int reopen = openat(procfd, "thread-self/fd/123",
                   &how, sizeof(how));
```

example: strict path safety



```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{ .flags = O_RDWR,
  .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };

int reopen = openat2(procfd, "thread-self/fd/123",
                     &how, sizeof(how));
```

example: strict path safety

```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{ .flags = O_DIRECTORY,
  .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int fd = openat2(procfd, "thread-self/fd", &how, sizeof(how));
/* validate no overmounts */
int reopen = openat(fd, "123", O_RDWR);
```

example: strict path safety

```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{ .flags = O_DIRECTORY,
  .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int fdffd = openat2(procfd, "thread-self/fd", &how, sizeof(how));
/* validate no overmounts */
int reopen = openat(fdffd, "123", O_RDWR);
```

example: strict path (un)safety



```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{ .flags = O_DIRECTORY,
  .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int fdffd = openat2(procfd, "thread-self/fd", &how, sizeof(how));
/* validate no overmounts */
int reopen = openat(fdffd, "123", O_RDWR);
```

man 2 fsopen

man 2 fspick

man 2 fsconfig

man 2 fsmount

man 2 open_tree*

man 2 move_mount



example: strict path safety

```
int procfd = open("/proc", O_DIRECTORY|O_PATH);
/* check PROC_SUPER_MAGIC and PROC_ROOT_INO */

struct open_how how =
{
    .flags = O_DIRECTORY,
    .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int fdffd = openat2(procfd, "thread-self/fd", &how, sizeof(how));
/* validate no overmounts */
int reopen = openat(fdffd, "123", O_RDWR);
```

example: strict path safety

```
int fsfd = fsopen("proc", 0);
fsconfig(fsfd, FSCONFIG_CMD_CREATE, NULL, 0, 0);
int procfd = fsmount(fsfd, 0, MOUNT_ATTR_NOEXEC|...);

struct open_how how =
{
    .flags = O_DIRECTORY,
    .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };
int fdffd = openat2(procfd, "thread-self/fd", &how, sizeof(how));
/* for open_tree(2) -- validate no overmounts */
int reopen = openat(fdffd, "123", O_RDWR);
```

example: strict path safety

```
int fsfd = fsopen("proc", 0);
fsconfig(fsfd, FSCONFIG_CMD_CREATE, NULL, 0, 0);
int procfd = fsmount(fsfd, 0, MOUNT_ATTR_NOEXEC|...);

struct open_how how =
{
    .flags = O_DIRECTORY,
    .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };

int fdffd = openat2(procfd, "thread-self/fd", &how, sizeof(how));
/* for open_tree(2) -- validate no overmounts */
int reopen = openat(fdffd, "123", O_RDWR);
```

example: strict path safety

```
int fsfd = fsopen("proc", 0);
fsconfig(fsfd, FSCONFIG_CMD_CREATE, NULL, 0, 0);
int procfd = fsmount(fsfd, 0, MOUNT_ATTR_NOEXEC|...);

struct open_how how =
{
    .flags = O_DIRECTORY,
    .resolve = RESOLVE_BENEATH|RESOLVE_NO_XDEV };

int fdffd = openat2(procfd, "thread-self/fd", &how, sizeof(how));
/* for open_tree(2) -- validate no overmounts */
int reopen = openat(fdffd, "123", O_RDWR);
```

example: libpathrs (c)

```
#include <pathrs.h>

int lfd = pathrs_proc_open(PATHRS_PROC_SELF,
                            "attr/exec", O_WRONLY);
dprintf(lfd, "exec docker-default\n");

int pfd = pathrs_proc_open(PATHRS_PROC_ROOT,
                            "sys/net/ipv4/ping_group_range",
                            O_RDWR);
dprintf(pfd, "0 0\n");

int reopen = pathrs_reopen(123, O_RDWR); // or pathrs_proc_open
```

example: libpathrs (rust)

```
use pathrs::{flags::OpenFlags, procfs::*};  
  
// Open *regular* file.  
let attr_file = ProcfsHandle::new()?  
    .open(  
        ProcfsBase::ProcThreadSelf,  
        "attr/exec",  
        OpenFlags::O_WRONLY,  
    )?;  
  
// Create your own private handle.  
let handle =  
    ProcfsHandleBuilder::new()  
        .unmasked(true)  
        .build()?;
  
// Open a magic-link.  
let exe = ProcfsHandle::new()?.open_follow(  
    ProcfsBase::ProcSelf,  
    "exe",  
    OpenFlags::O_RDONLY,  
)?;  
  
// Equivalent to readlinkat(fd, "").  
let fd_path = ProcfsHandle::new()?.readlink(  
    ProcfsBase::ProcThreadSelf,  
    format!("fd/{}", file.as_raw_fd()),  
    OpenFlags::O_RDONLY,  
)?;
```

—

“wait, why do we
care about this?”



runc

threat models

- runc (currently) has no real threat model.
 - Most vulnerabilities have been “misconfiguration” bugs.
 - Higher-level runtimes let unprivileged users do wacky things.
 - Most people **still** don’t use user namespaces.
-

—
pop quiz!

is this a vulnerability? (i)

```
$ cat config.json
{
    "linux": { "namespaces": [] },
    ...
}
```

is this a vulnerability? (i)

```
$ cat config.json
{
    "linux": { "namespaces": [] },
    ...
}
```



is this a vulnerability? (ii)

```
$ cat config.json
{
    "mounts": [
        ...,
        { "destination": "/host",
          "source": "/",
          "type": "bind",
          "options": [ "rbind" ] }
    ],
    ...
}
```

is this a vulnerability? (ii)

```
$ cat config.json
{
    "mounts": [
        ...,
        { "destination": "/host",
          "source": "/",
          "type": "bind",
          "options": [ "rbind" ] }
    ],
    ...
}
```



is this a vulnerability? (iii)

```
$ cat config.json
{
    "mounts": [
        ...,
        {
            "destination": "/volume",
            "source": "/some/volume/path",
            "type": "bind",
            "options": [ "rbind" ] },
        ...
    ],
    ...
}
```

is this a vulnerability? (iii)

```
$ cat config.json
{
    "mounts": [
        ...,
        {
            "destination": "/volume",
            "source": "/some/volume/path",
            "type": "bind",
            "options": [ "rbind" ] },
        ...
    ],
    ...
}
```



CVE-2025-31133 / CVE-2025-52565

```
$ cat config.json
{
    "mounts": [
        ...,
        {
            "destination": "/volume",
            "source": "/some/volume/path",
            "type": "bind",
            "options": [ "rbind" ] },
        ...
    ],
    ...
}
```



CVE-2025-31133 / CVE-2025-52565

- `/volume` is a symlink to `/dev`.
- Racing process swaps files in `/some/volume/path` with symlink to `/proc/sys/kernel/core_pattern`:
 - CVE-2025-31133 – `/dev/null` (masked files)
 - CVE-2025-52565 – `/dev/pts/0` (`/dev/console`)
- Bind-mount *source* becomes `/proc/sys/kernel/core_pattern`, creating a `rw` bind-mount to a masked procfs file.

solutions

- Added much stricter validation of special inodes we use.
 - *Takeaway:* “Safe” major:minor numbers are very handy.
- Mountpoint creation was moved to [libpathrs \(pathrs-lite\)](#).
 - Everything is now (mostly) file-descriptor-based.
 - *Takeaway:* Wow, regular path safety is a good idea!
- [TIOCGPTPEER](#) for consoles.

is this a vulnerability? (iv)

```
$ cat config.json
{
    "mounts": [
        ...,
        {
            "destination": "/proc",
            "source": "/fake/procfs",
            "type": "bind",
            "options": [ "rbind" ] },
        ],
        ...
}
```

is this a vulnerability? (iv)

```
$ cat config.json
{
    "mounts": [
        ...,
        {
            "destination": "/proc",
            "source": "/fake/procfs",
            "type": "bind",
            "options": [ "rbind" ] },
        ],
        ...
}
```



is this a vulnerability? (v)

```
$ cat config.json
{
    "mounts": [
        ...,
        {
            "destination": "/proc/1/attr/apparmor/exec",
            "source": "/proc/1/sched",
            "type": "bind",
            "options": [ "rbind" ] },
        ],
        ...
    }
```

is this a vulnerability? (v)



```
$ cat config.json
{
  "mounts": [
    ...,
    {
      "destination": "/proc/1/attr/apparmor/exec",
      "source": "/proc/1/sched",
      "type": "bind",
      "options": [ "rbind" ] },
    ],
    ...
}
```

is this a vulnerability? (v)



```
$ cat config.json
{
  "mounts": [
    ...,
    {
      "destination": "/proc/1/attr/apparmor/exec",
      "source": "/proc/1/sched",
      "type": "bind",
      "options": [ "rbind" ] },
    ],
    ...
}
```



is this a vulnerability? (vi)

```
$ cat config.json
{
    "mounts": [
        ...,
        { "destination": "/foo", "source": "/some/cache/path",
          "type": "bind", "options": [ "rbind" ] },
        { "destination": "/foo/link/thread-self/attr/apparmor",
          "source": "/some/other-cache/path",
          "type": "bind", "options": [ "rbind" ] },
    ],
    ...
}
```

is this a vulnerability? (vi)



```
$ cat config.json
{
  "mounts": [
    ...,
    { "destination": "/foo", "source": "/some/cache/path",
      "type": "bind", "options": [ "rbind" ] },
    { "destination": "/foo/link/thread-self/attr/apparmor",
      "source": "/some/other-cache/path",
      "type": "bind", "options": [ "rbind" ] },
  ],
  ...
}
```

CVE-2025-52881



```
$ cat config.json
{
  "mounts": [
    ...,
    { "destination": "/foo", "source": "/tmp/cache-foo",
      "type": "bind", "options": [ "rbind" ] },
    { "destination": "/foo/link/thread-self/attr/apparmor",
      "source": "/tmp/cache-apparmor",
      "type": "bind", "options": [ "rbind" ] },
  ],
  ...
}
```

CVE-2025-52881

- `/tmp/cache-apparmor/exec` is a symlink to a fun target:
 - `/proc/1/sched` – (no-op)
 - `/proc/sysrq-trigger` (crash – "exec docker-default")
- Racing process swaps `/foo/link` symlink between `/proc` and dummy directory.
 - This bypassed our anti-`/proc` mount checks.
- `/proc/sys/kernel/core_pattern` allows a container escape.

solutions

- Additional hardening when doing mounts.
 - *Takeaway:* We really should've switched to the new mount API much earlier.
- Switch to [libpathrs](#) ([pathrs-lite](#)) procfs API for writes.
 - Also, audited all write paths for misdirectable writes.
 - *Takeaway:* Protecting against everything is quite hard.

runc todos

- Move our mount infrastructure to `fsopen` / `open_tree`.
 - Masked path application will finally be atomic.
- Do a deeper audit of all path-based code in runc.
- Switch to `libpathrs` for runc builds.
 - `pathrs-lite` can use `libpathrs` as a backend.

kernel todos

- Still some nice-to-have extensions for openat2.
 - RESOLVE_NO_DOTDOT?
 - rrootfd / cwdfd split?
- Blocking all magic-link overmounts would help a lot.
 - Most have been blocked since 6.12.

general takeaways

- Use `openat2` or `libpathrs` (for Go, maybe `pathrs-lite`).
 - Switch to a more file-descriptor based design.
 - Doubly so if you need to touch `/proc`.
 - *Every pathname syscall is potentially dangerous.*
-

ΛMUTΛBLE

questions?

(rants, pitchforks...?)

CC-BY-SA 4.0



fun issues

apparmor, d_path, pain

⋮  CVE-2025-52881: fd reopening ...
causes issues with AppArmor profiles
(open sysctl
net.ipv4.ip_unprivileged_port_start file:
reopen fd 8: permission denied)
#4968 · cyphar opened on Nov 6  92

apparmor, d_path, pain

- Nested containers (under LXC).
- AppArmor returns -EACCES for `/proc/sys/net/...` write.
- Triggered by switch to `fsopen("proc")`.

apparmor, d_path, pain

```
deny /proc/sys/[^fknu]*{,/**} wklx,  
deny /proc/sys/n[^e]*{,/**} wklx,  
deny /proc/sys/ne[^t]*{,/**} wklx,  
deny /proc/sys/net?*{,/**} wklx,  
...
```

apparmor, d_path, pain

```
deny /proc/sys/[^fknu]*{,/**} wklx,  
deny /proc/sys/n[^e]*{,/**} wklx,  
deny /proc/sys/ne[^t]*{,/**} wklx,  
deny /proc/sys/net?*{,/**} wklx,  
...  
  
deny /sys/[^fdck]*{,/**} wklx,  
deny /sys/k[^e]*{,/**} wklx,  
...
```

apparmor, d_path, pain

```
deny /proc/sys/[^fknu]*{,/**} wklx,  
deny /proc/sys/n[^e]*{,/**} wklx,  
deny /proc/sys/ne[^t]*{,/**} wklx,  
deny /proc/sys/net?*{,/**} wklx,
```

...

```
deny /sys/[^fdck]*{,/**} wklx,  
deny /sys/k[^e]*{,/**} wklx,
```

...

runc **AppArmor**

func AppArmor

```
int procfd = fsopen("proc");    aa_path_name() ⇒ "/"
```

func AppArmor

```
int procfd = fsopen("proc");
```

```
aa_path_name() ⇒ "/"
```

```
int fd = openat(procfd,  
                 "sys/foo/bar", ...);
```

```
aa_path_name() ⇒ "/sys/foo/bar"
```

func AppArmor

```
int procfd = fsopen("proc");
```

aa_path_name() ⇒ "/"

```
int fd = openat(procfd,  
                 "sys/foo/bar", ...);
```

aa_path_name() ⇒ "/sys/foo/bar"



apparmor, d_path, pain

```
int fsfd = fsopen("tmpfs", 0);
fsconfig(fsfd, FSCONFIG_CMD_CREATE, NULL, 0, 0);
int tmpfd = fsmount(fsfd, 0, MOUNT_ATTR_NOEXEC|...);

mkdirat(tmpfd, "sys", 0755);
openat(tmpfd, "sys/foo", O_CREAT, 0755); /* EACCES */
```

apparmor, d_path, pain

incusd/apparmor/lxc: Don't bother with sys/proc protections
when nesting enabled #2624



stgraber merged 1 commit into `lxc:main` from `stgraber:main` on Nov 6

dangling symlinks

- Previously we would expand dangling symlinks.
 - `/foo/bar` → `/some/non-existant/path`
 - `/foo/bar/baz` → `/some/non-existant/path/baz`
- Some users depend on this behaviour...
 - Emulating it with file descriptors is quite hard.
- *Solution:* Expand the path and then use that as the user path.

ΛMUTΛBLE

questions?

(rants, pitchforks...?)

CC-BY-SA 4.0

